

**Implement K-Means clustering/  
hierarchical clustering on  
sales\_data\_sample.csv dataset. Determine  
the number of clusters using the elbow  
method.**

```
In [1]: import pandas as pd  
import numpy as np
```

```
In [2]: df = pd.read_csv('7458_sales_data_sample.csv', encoding='unicode_escape')
```

```
In [3]: df.head
```

Out[3]: <bound method NDFrame.head of LINENUMBER SALES \		ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDER
0	10107	30	95.70	2	2871.00
1	10121	34	81.35	5	2765.90
2	10134	41	94.74	2	3884.34
3	10145	45	83.26	6	3746.70
4	10159	49	100.00	14	5205.27
...	...	...	...	...	...
2818	10350	20	100.00	15	2244.40
2819	10373	29	100.00	1	3978.51
2820	10386	43	100.00	4	5417.57
2821	10397	34	62.24	1	2116.16
2822	10414	47	65.52	9	3079.44
ORDERDATE STATUS QTR_ID MONTH_ID YEAR_ID ... \					
0	2/24/2003 0:00	Shipped	1	2	2003 ...
1	5/7/2003 0:00	Shipped	2	5	2003 ...
2	7/1/2003 0:00	Shipped	3	7	2003 ...
3	8/25/2003 0:00	Shipped	3	8	2003 ...
4	10/10/2003 0:00	Shipped	4	10	2003 ...
...	...	...	...	...	...
2818	12/2/2004 0:00	Shipped	4	12	2004 ...
2819	1/31/2005 0:00	Shipped	1	1	2005 ...
2820	3/1/2005 0:00	Resolved	1	3	2005 ...
2821	3/28/2005 0:00	Shipped	1	3	2005 ...
2822	5/6/2005 0:00	On Hold	2	5	2005 ...
ADDRESSLINE1 ADDRESSLINE2 CITY STATE \					
0	897 Long Airport Avenue	NaN	NYC	NY	
1	59 rue de l'Abbaye	NaN	Reims	NaN	
2	27 rue du Colonel Pierre Avia	NaN	Paris	NaN	
3	78934 Hillside Dr.	NaN	Pasadena	CA	
4	7734 Strong St.	NaN	San Francisco	CA	
...	...	...	...	...	...
2818	C/ Moralzarzal, 86	NaN	Madrid	NaN	
2819	Torikatu 38	NaN	Oulu	NaN	
2820	C/ Moralzarzal, 86	NaN	Madrid	NaN	
2821	1 rue Alsace-Lorraine	NaN	Toulouse	NaN	
2822	8616 Spinnaker Dr.	NaN	Boston	MA	
POSTALCODE COUNTRY TERRITORY CONTACTLASTNAME CONTACTFIRSTNAME DEALSIZE					
0	10022	USA	NaN	Yu	Kwai Small
1	51100	France	EMEA	Henriot	Paul Small
2	75508	France	EMEA	Da Cunha	Daniel Medium
3	90003	USA	NaN	Young	Julie Medium
4	NaN	USA	NaN	Brown	Julie Medium
...	...	...	...	...	...
2818	28034	Spain	EMEA	Freyre	Diego Small
2819	90110	Finland	EMEA	Koskitalo	Pirkko Medium
2820	28034	Spain	EMEA	Freyre	Diego Medium
2821	31000	France	EMEA	Roulet	Annette Small
2822	51003	USA	NaN	Yoshido	Juri Medium

[2823 rows x 25 columns]&gt;

```
In [4]: df.info
```

Out[4]: <bound method DataFrame.info of

	ERLINENUMBER	SALES	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORD
0	10107		30	95.70	2	2871.00
1	10121		34	81.35	5	2765.90
2	10134		41	94.74	2	3884.34
3	10145		45	83.26	6	3746.70
4	10159		49	100.00	14	5205.27
...	...	...	...	...	...	...
2818	10350		20	100.00	15	2244.40
2819	10373		29	100.00	1	3978.51
2820	10386		43	100.00	4	5417.57
2821	10397		34	62.24	1	2116.16
2822	10414		47	65.52	9	3079.44

	ORDERDATE	STATUS	QTR_ID	MONTH_ID	YEAR_ID	...	\
0	2/24/2003 0:00	Shipped	1	2	2003	...	
1	5/7/2003 0:00	Shipped	2	5	2003	...	
2	7/1/2003 0:00	Shipped	3	7	2003	...	
3	8/25/2003 0:00	Shipped	3	8	2003	...	
4	10/10/2003 0:00	Shipped	4	10	2003	...	
...	...	...	...	...	...	...	...
2818	12/2/2004 0:00	Shipped	4	12	2004	...	
2819	1/31/2005 0:00	Shipped	1	1	2005	...	
2820	3/1/2005 0:00	Resolved	1	3	2005	...	
2821	3/28/2005 0:00	Shipped	1	3	2005	...	
2822	5/6/2005 0:00	On Hold	2	5	2005	...	

	ADDRESSLINE1	ADDRESSLINE2	CITY	STATE	\
0	897 Long Airport Avenue	NaN	NYC	NY	
1	59 rue de l'Abbaye	NaN	Reims	NaN	
2	27 rue du Colonel Pierre Avia	NaN	Paris	NaN	
3	78934 Hillside Dr.	NaN	Pasadena	CA	
4	7734 Strong St.	NaN	San Francisco	CA	
...	...	...	...	...	...
2818	C/ Moralzarzal, 86	NaN	Madrid	NaN	
2819	Torikatu 38	NaN	Oulu	NaN	
2820	C/ Moralzarzal, 86	NaN	Madrid	NaN	
2821	1 rue Alsace-Lorraine	NaN	Toulouse	NaN	
2822	8616 Spinnaker Dr.	NaN	Boston	MA	

	POSTALCODE	COUNTRY	TERRITORY	CONTACTLASTNAME	CONTACTFIRSTNAME	DEALSIZE
0	10022	USA	NaN	Yu	Kwai	Small
1	51100	France	EMEA	Henriot	Paul	Small
2	75508	France	EMEA	Da Cunha	Daniel	Medium
3	90003	USA	NaN	Young	Julie	Medium
4	NaN	USA	NaN	Brown	Julie	Medium
...	...	...	...	...	...	...
2818	28034	Spain	EMEA	Freyre	Diego	Small
2819	90110	Finland	EMEA	Koskitalo	Pirkko	Medium
2820	28034	Spain	EMEA	Freyre	Diego	Medium
2821	31000	France	EMEA	Roulet	Annette	Small
2822	51003	USA	NaN	Yoshido	Juri	Medium

[2823 rows x 25 columns]>

```
In [5]: #Columns to Remove  
to_drop = ['ADDRESSLINE1', 'ADDRESSLINE2', 'STATE', 'POSTALCODE', 'PHONE']  
df = df.drop(to_drop, axis=1)
```

```
In [6]: #Check for null values  
df.isnull().sum()
```

```
Out[6]: ORDERNUMBER      0  
QUANTITYORDERED      0  
PRICEEACH            0  
ORDERLINENUMBER      0  
SALES                0  
ORDERDATE             0  
STATUS                0  
QTR_ID                0  
MONTH_ID               0  
YEAR_ID                0  
PRODUCTLINE            0  
MSRP                  0  
PRODUCTCODE            0  
CUSTOMERNAME           0  
CITY                  0  
COUNTRY                0  
TERRITORY              1074  
CONTACTLASTNAME        0  
CONTACTFIRSTNAME       0  
DEALSIZE                0  
dtype: int64
```

```
In [7]: #Bhai bhai Look at territory  
#But territory does not have significant impact on analysis, let it be
```

```
In [8]: df.dtypes
```

```
Out[8]: ORDERNUMBER      int64  
QUANTITYORDERED      int64  
PRICEEACH            float64  
ORDERLINENUMBER      int64  
SALES                float64  
ORDERDATE             object  
STATUS                object  
QTR_ID                int64  
MONTH_ID               int64  
YEAR_ID                int64  
PRODUCTLINE            object  
MSRP                  int64  
PRODUCTCODE            object  
CUSTOMERNAME           object  
CITY                  object  
COUNTRY                object  
TERRITORY              object  
CONTACTLASTNAME        object  
CONTACTFIRSTNAME       object  
DEALSIZE                object  
dtype: object
```

```
In [9]: #ORDERDATE Should be in date time
df['ORDERDATE'] = pd.to_datetime(df['ORDERDATE'])
```

```
In [10]: #We need to create some features in order to create cluseters
#Recency: Number of days between customer's latest order and today's date
#Frequency : Number of purchases by the customers
#MonetaryValue : Revenue generated by the customers
import datetime as dt
snapshot_date = df['ORDERDATE'].max() + dt.timedelta(days = 1)
df_RFMs = df.groupby(['CUSTOMERNAME']).agg({
    'ORDERDATE' : lambda x : (snapshot_date - x.max()).days,
    'ORDERNUMBER' : 'count',
    'SALES' : 'sum'
})

#Rename the columns
df_RFMs.rename(columns = {
    'ORDERDATE' : 'Recency',
    'ORDERNUMBER' : 'Frequency',
    'SALES' : 'MonetaryValue'
}, inplace=True)
```

```
In [11]: df_RFMs.head()
```

Out[11]:

CUSTOMERNAME	Recency	Frequency	MonetaryValue
<b>AV Stores, Co.</b>	196	51	157807.81
<b>Alpha Cognac</b>	65	20	70488.44
<b>Amica Models &amp; Co.</b>	265	26	94117.26
<b>Anna's Decorations, Ltd</b>	84	46	153996.13
<b>Atelier graphique</b>	188	7	24179.96

```
In [12]: # Divide into segments
# We create 4 quartile ranges
df_RFMs['M'] = pd.qcut(df_RFMs['MonetaryValue'], q = 4, labels = range(1,5))
df_RFMs['R'] = pd.qcut(df_RFMs['Recency'], q = 4, labels = list(range(4,0,-1)))
df_RFMs['F'] = pd.qcut(df_RFMs['Frequency'], q = 4, labels = range(1,5))

df_RFMs.head()
```

Out[12]:

	Recency	Frequency	MonetaryValue	M	R	F
<b>CUSTOMERNAME</b>						

<b>AV Stores, Co.</b>	196	51	157807.81	4	2	4
<b>Alpha Cognac</b>	65	20	70488.44	2	4	2
<b>Amica Models &amp; Co.</b>	265	26	94117.26	3	1	2
<b>Anna's Decorations, Ltd</b>	84	46	153996.13	4	3	4
<b>Atelier graphique</b>	188	7	24179.96	1	2	1

In [13]:

```
#Create another column for RFM score
df_RFM['RFM_Score'] = df_RFM[['R', 'M', 'F']].sum(axis=1)
df_RFM.head()
```

Out[13]:

	Recency	Frequency	MonetaryValue	M	R	F	RFM_Score
<b>CUSTOMERNAME</b>							

<b>AV Stores, Co.</b>	196	51	157807.81	4	2	4	10
<b>Alpha Cognac</b>	65	20	70488.44	2	4	2	8
<b>Amica Models &amp; Co.</b>	265	26	94117.26	3	1	2	6
<b>Anna's Decorations, Ltd</b>	84	46	153996.13	4	3	4	11
<b>Atelier graphique</b>	188	7	24179.96	1	2	1	4

## We create levels for our Customers

RFM Score > 10 : High Value Customers

RFM Score < 10 and RFM Score >= 6 : Mid Value Customers

RFM Score < 6 : Low Value Customers

In [14]:

```
def rfm_level(df):
    if bool(df['RFM_Score'] >= 10):
        return 'High Value Customer'

    elif bool(df['RFM_Score'] < 10) and bool(df['RFM_Score'] >= 6):
        return 'Mid Value Customer'
    else:
        return 'Low Value Customer'

df_RFM['RFM_Level'] = df_RFM.apply(rfm_level, axis = 1)
df_RFM.head()
```

Out[14]:

CUSTOMERNAME	Recency	Frequency	MonetaryValue	M	R	F	RFM_Score	RFM_Level
<b>AV Stores, Co.</b>	196	51	157807.81	4	2	4	10	High Value Customer
<b>Alpha Cognac</b>	65	20	70488.44	2	4	2	8	Mid Value Customer
<b>Amica Models &amp; Co.</b>	265	26	94117.26	3	1	2	6	Mid Value Customer
<b>Anna's Decorations, Ltd</b>	84	46	153996.13	4	3	4	11	High Value Customer
<b>Atelier graphique</b>	188	7	24179.96	1	2	1	4	Low Value Customer

In [15]:

```
# Time to perform KMeans
data = df_RFMs[['Recency', 'Frequency', 'MonetaryValue']]
data.head()
```

Out[15]:

CUSTOMERNAME	Recency	Frequency	MonetaryValue
<b>AV Stores, Co.</b>	196	51	157807.81
<b>Alpha Cognac</b>	65	20	70488.44
<b>Amica Models &amp; Co.</b>	265	26	94117.26
<b>Anna's Decorations, Ltd</b>	84	46	153996.13
<b>Atelier graphique</b>	188	7	24179.96

In [16]:

```
# Our data is skewed we must remove it by performing Log transformation
data_log = np.log(data)
data_log.head()
```

Out[16]:

CUSTOMERNAME	Recency	Frequency	MonetaryValue
<b>AV Stores, Co.</b>	5.278115	3.931826	11.969133
<b>Alpha Cognac</b>	4.174387	2.995732	11.163204
<b>Amica Models &amp; Co.</b>	5.579730	3.258097	11.452297
<b>Anna's Decorations, Ltd</b>	4.430817	3.828641	11.944683
<b>Atelier graphique</b>	5.236442	1.945910	10.093279

```
In [17]: #Standardization
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(data_log)
data_normalized = scaler.transform(data_log)
data_normalized = pd.DataFrame(data_normalized, index = data_log.index, columns=dat
data_normalized.describe().round(2)
```

Out[17]:

	Recency	Frequency	MonetaryValue
<b>count</b>	92.00	92.00	92.00
<b>mean</b>	0.00	-0.00	0.00
<b>std</b>	1.01	1.01	1.01
<b>min</b>	-3.51	-3.67	-3.82
<b>25%</b>	-0.24	-0.41	-0.39
<b>50%</b>	0.37	0.06	-0.04
<b>75%</b>	0.53	0.45	0.52
<b>max</b>	1.12	4.03	3.92

```
In [18]: #Fit KMeans and use elbow method to choose the number of clusters
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.cluster import KMeans

sse = {}

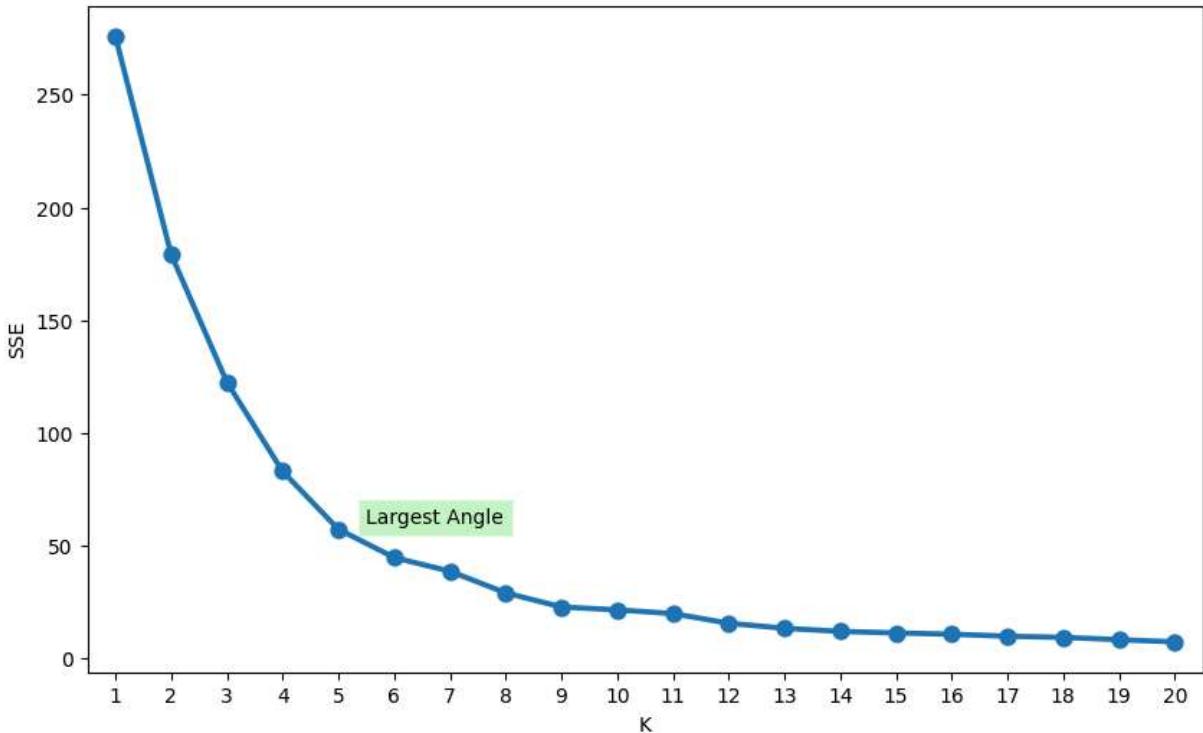
for k in range(1, 21):
    kmeans = KMeans(n_clusters = k, random_state = 1)
    kmeans.fit(data_normalized)
    sse[k] = kmeans.inertia_
```

```
In [19]: plt.figure(figsize=(10,6))
plt.title('The Elbow Method')

plt.xlabel('K')
plt.ylabel('SSE')
plt.style.use('ggplot')

sns.pointplot(x=list(sse.keys()), y = list(sse.values()))
plt.text(4.5, 60, "Largest Angle", bbox = dict(facecolor = 'lightgreen', alpha = 0.5))
plt.show()
```

## The Elbow Method



```
In [20]: # 5 number of clusters seems good
kmeans = KMeans(n_clusters=5, random_state=1)
kmeans.fit(data_normalized)
cluster_labels = kmeans.labels_

data_rfkm = data.assign(Cluster = cluster_labels)
data_rfkm.head()
```

Out[20]:

	Recency	Frequency	MonetaryValue	Cluster
--	---------	-----------	---------------	---------

**CUSTOMERNAME**

<b>AV Stores, Co.</b>	196	51	157807.81	4
<b>Alpha Cognac</b>	65	20	70488.44	2
<b>Amica Models &amp; Co.</b>	265	26	94117.26	2
<b>Anna's Decorations, Ltd</b>	84	46	153996.13	4
<b>Atelier graphique</b>	188	7	24179.96	1